

Objective: To provide hands-on experience to students in implementing and applying linear regression models and evaluate the performance.

1. Download the following dataset: Wine Quality Dataset: Predicting wine quality based on various features. URL:

<https://archive.ics.uci.edu/dataset/186/wine+quality>

```
1 import pandas as pd
2 from sklearn.model_selection import train_test_split
3 from sklearn.linear_model import LinearRegression
4 from sklearn.preprocessing import PolynomialFeatures
5 from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
6
1 url = "https://archive.ics.uci.edu/ml/machine-learning-databases/wine-quality/winequality-white.csv"
2 df = pd.read_csv(url, sep=";")
3
1 print("Statistical values of the dataset:")
2 print(df.describe())
3 print("\nShape of the data:", df.shape)
```



Statistical values of the dataset:

	fixed acidity	volatile acidity	citric acid	residual sugar \
count	4898.000000	4898.000000	4898.000000	4898.000000
mean	6.854788	0.278241	0.334192	6.391415
std	0.843868	0.100795	0.121020	5.072058
min	3.800000	0.080000	0.000000	0.600000
25%	6.300000	0.210000	0.270000	1.700000
50%	6.800000	0.260000	0.320000	5.200000
75%	7.300000	0.320000	0.390000	9.900000
max	14.200000	1.100000	1.660000	65.800000

	chlorides	free sulfur dioxide	total sulfur dioxide	density \
count	4898.000000	4898.000000	4898.000000	4898.000000
mean	0.045772	35.308085	138.360657	0.994027
std	0.021848	17.007137	42.498065	0.002991
min	0.009000	2.000000	9.000000	0.987110
25%	0.036000	23.000000	108.000000	0.991723
50%	0.043000	34.000000	134.000000	0.993740
75%	0.050000	46.000000	167.000000	0.996100
max	0.346000	289.000000	440.000000	1.038980

	pH	sulphates	alcohol	quality
count	4898.000000	4898.000000	4898.000000	4898.000000
mean	3.188267	0.489847	10.514267	5.877909
std	0.151001	0.114126	1.230621	0.885639
min	2.720000	0.220000	8.000000	3.000000

25%	3.090000	0.410000	9.500000	5.000000
50%	3.180000	0.470000	10.400000	6.000000
75%	3.280000	0.550000	11.400000	6.000000
max	3.820000	1.080000	14.200000	9.000000

Shape of the data: (4898, 12)

```
1 X = df.drop('quality', axis=1)
2 Y = df['quality']
```

(4898, 11)

```
1 print("\nShape of X (features):", X.shape)
2 print("Shape of Y (target variable):", Y.shape)
```

Shape of X (features): (4898, 11)
Shape of Y (target variable): (4898,)

```
1
2 df_train, df_test = train_test_split(df, test_size=0.2, random_state=42)
3
4
```

```
1 lr_model = LinearRegression()
2 lr_model.fit(df_train.drop('quality', axis=1), df_train['quality'])
```

▼ LinearRegression

LinearRegression()

```
1 # Polynomial Regression model
2 poly = PolynomialFeatures(degree=2)
3 X_poly = poly.fit_transform(df_train.drop('quality', axis=1))
4 poly_model = LinearRegression()
5 poly_model.fit(X_poly, df_train['quality'])
6
```

▼ LinearRegression

LinearRegression()

```
1 def evaluate_model(model, X, y_true):
2     y_pred = model.predict(X)
3     mse = mean_squared_error(y_true, y_pred)
```

```
4 rmse = mean_squared_error(y_true, y_pred, squared=False)
5 mae = mean_absolute_error(y_true, y_pred)
6 r2 = r2_score(y_true, y_pred)
7 return mse, rmse, mae, r2

1 mse_lr, rmse_lr, mae_lr, r2_lr = evaluate_model(lr_model, df_test.drop('quality', axis=1), df_test['quality'])
2 print("\nLinear Regression Model Evaluation:")
3 print("MSE:", mse_lr)
4 print("RMSE:", rmse_lr)
5 print("MAE:", mae_lr)
6 print("R2 Score:", r2_lr)
7
```

```
Linear Regression Model Evaluation:
MSE: 0.5690247717229278
RMSE: 0.754337306331145
MAE: 0.5862665383250473
R2 Score: 0.2652750042179125
```

```
1 X_poly_test = poly.transform(df_test.drop('quality', axis=1))
2 mse_poly, rmse_poly, mae_poly, r2_poly = evaluate_model(poly_model, X_poly_test, df_test['quality'])
3 print("\nPolynomial Regression Model Evaluation:")
4 print("MSE:", mse_poly)
5 print("RMSE:", rmse_poly)
6 print("MAE:", mae_poly)
7 print("R2 Score:", r2_poly)
```

```
Polynomial Regression Model Evaluation:
MSE: 0.6193710931852203
RMSE: 0.7870013298497153
MAE: 0.5669837736822492
R2 Score: 0.20026781531816773
```

IRIS DATASET

```
1 import pandas as pd
2
3
4 column_names = ["sepal_length", "sepal_width", "petal_length", "petal_width", "class"]
5 df_iris = pd.read_csv("https://archive.ics.uci.edu/ml/machine-learning-databases/iris/iris.data", header=None, names=column_names)
6
7 print(df_iris.describe())
```

```
8 print(df_iris.shape)
```

```

      count  sepal_length  sepal_width  petal_length  petal_width
mean      5.843333      3.054000      3.758667      1.198667
std        0.828066      0.433594      1.764420      0.763161
min         4.300000      2.000000      1.000000      0.100000
25%         5.100000      2.800000      1.600000      0.300000
50%         5.800000      3.000000      4.350000      1.300000
75%         6.400000      3.300000      5.100000      1.800000
max         7.900000      4.400000      6.900000      2.500000
(150, 5)
```

```
1 from sklearn.model_selection import train_test_split
```

```
2
```

```
3 X_iris = df_iris.drop('class', axis=1)
```

```
4 Y_iris = df_iris['class']
```

```
5
```

```
6 X_iris_train, X_iris_test, Y_iris_train, Y_iris_test = train_test_split(X_iris, Y_iris, test_size=0.2, random_state=42)
```

```
7
```

```
1 from sklearn.linear_model import LogisticRegression
```

```
2
```

```
3 lr_iris = LogisticRegression(max_iter=200).fit(X_iris_train, Y_iris_train)
```

```
4 predictions_iris = lr_iris.predict(X_iris_test)
```

```
5
```

```
1 from sklearn.metrics import accuracy_score, precision_recall_fscore_support, roc_auc_score, roc_curve, auc
```

```
2 import numpy as np
```

```
3
```

```
4 accuracy = accuracy_score(Y_iris_test, predictions_iris)
```

```
5 precision, recall, fscore, _ = precision_recall_fscore_support(Y_iris_test, predictions_iris, average='macro')
```

```
6
```

```
7 from sklearn.preprocessing import label_binarize
```

```
8 from sklearn.multiclass import OneVsRestClassifier
```

```
9
```

```
10 Y_iris_test_bin = label_binarize(Y_iris_test, classes=["Iris-setosa", "Iris-versicolor", "Iris-virginica"])
```

```
11 n_classes = Y_iris_test_bin.shape[1]
```

```
12
```

```
13 classifier = OneVsRestClassifier(LogisticRegression(max_iter=200))
```

```
14 score = classifier.fit(X_iris_train, label_binarize(Y_iris_train, classes=["Iris-setosa", "Iris-versicolor", "Iris-virginica"])).decision_function(X_iris_test)
```

```
15
```

```
16 fpr = dict()
```

```
17 tpr = dict()
```

```
18 roc_auc = dict()
```

```
19
```

```
20 for i in range(n_classes):
```

```
21 fpr[i], tpr[i], _ = roc_curve(Y_iris_test_bin[:, i], score[:, i])
22 roc_auc[i] = auc(fpr[i], tpr[i])
23
24 fpr["micro"], tpr["micro"], _ = roc_curve(Y_iris_test_bin.ravel(), score.ravel())
25 roc_auc["micro"] = auc(fpr["micro"], tpr["micro"])
26
27 print("Accuracy:", accuracy)
28 print("Precision:", precision)
29 print("Recall:", recall)
30 print("F-score:", fscore)
31 print("AUC (micro-averaged):", roc_auc["micro"])
32
Accuracy: 1.0
Precision: 1.0
Recall: 1.0
F-score: 1.0
AUC (micro-averaged): 0.9794444444444446
```

Colab paid products - Cancel contracts here

✓ 0s completed at 12:01

